

# 2-Pandas

September 3, 2016

## 1 Getting to know your data with Pandas

### 1.1 Pandas

Pandas is the Python Data Analysis Library.

Pandas is an extremely versatile tool for manipulating datasets.

It also produces high quality plots with matplotlib, and integrates nicely with other libraries that expect NumPy arrays.

The most important tool provided by Pandas is the **data frame**.

A data frame is a table in which each row and column is given a label.

### 1.2 Getting started

### 1.3 Fetching, storing and retrieving your data

More information on what types of data you can fetch

[http://pandas.pydata.org/pandas-docs/stable/remote\\_data.html](http://pandas.pydata.org/pandas-docs/stable/remote_data.html)

```
In [109]: stocks = 'YELP'
          data_source = 'yahoo'
          start = datetime(2015,1,1)
          end = datetime(2015,12,31)

          yahoo_stocks = web.DataReader(stocks, data_source, start, end)

          # yahoo_stocks.head()
          yahoo_stocks.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 252 entries, 2015-01-02 to 2015-12-31
Data columns (total 6 columns):
Open                252 non-null float64
High                252 non-null float64
Low                 252 non-null float64
Close              252 non-null float64
Volume              252 non-null int64
Adj Close           252 non-null float64
dtypes: float64(5), int64(1)
```

memory usage: 13.8 KB

In [110]: yahoo\_stocks

```
Out[110]:
```

	Date	Open	High	Low	Close	Volume	Adj Clo
	2015-01-02	55.459999	55.599998	54.240002	55.150002	1664500	55.1500
	2015-01-05	54.540001	54.950001	52.330002	52.529999	2023000	52.5299
	2015-01-06	52.549999	53.930000	50.750000	52.439999	3762800	52.4399
	2015-01-07	53.320000	53.750000	51.759998	52.209999	1548200	52.2099
	2015-01-08	52.590000	54.139999	51.759998	53.830002	2015300	53.8300
	2015-01-09	55.959999	56.990002	54.720001	56.070000	6222600	56.0700
	2015-01-12	56.000000	56.060001	53.430000	54.020000	2405100	54.0200
	2015-01-13	54.470001	54.799999	52.520000	53.180000	1952100	53.1800
	2015-01-14	52.799999	53.680000	51.459999	52.200001	1854600	52.2000
	2015-01-15	53.000000	53.610001	50.029999	50.119999	2640400	50.1199
	2015-01-16	50.180000	51.490002	50.029999	51.389999	2183300	51.3899
	2015-01-20	51.650002	51.779999	50.689999	51.410000	1227600	51.4100
	2015-01-21	51.200001	53.500000	51.200001	53.410000	3248100	53.4100
	2015-01-22	53.869999	55.279999	53.119999	54.799999	2295400	54.7999
	2015-01-23	54.660000	55.639999	54.299999	55.189999	1636400	55.1899
	2015-01-26	55.119999	55.790001	54.830002	55.410000	1450300	55.4100
	2015-01-27	56.060001	56.160000	54.570000	55.630001	2410400	55.6300
	2015-01-28	56.150002	56.150002	52.919998	53.000000	2013100	53.0000
	2015-01-29	52.849998	53.310001	51.410000	52.930000	1844100	52.9300
	2015-01-30	52.590000	53.419998	52.049999	52.470001	1875400	52.4700
	2015-02-02	52.939999	53.500000	51.209999	53.470001	2105500	53.4700
	2015-02-03	53.830002	55.930000	53.410000	55.779999	2885400	55.7799
	2015-02-04	55.529999	57.070000	55.250000	56.740002	2498600	56.7400
	2015-02-05	57.599998	57.700001	56.080002	57.470001	4657300	57.4700
	2015-02-06	47.700001	48.169998	44.860001	45.110001	25137400	45.1100
	2015-02-09	44.910000	45.040001	42.099998	42.169998	13079300	42.1699
	2015-02-10	43.830002	45.549999	43.310001	44.660000	11267700	44.6600
	2015-02-11	45.389999	46.430000	44.810001	46.180000	6359400	46.1800
	2015-02-12	46.450001	47.840000	45.950001	47.630001	4375000	47.6300
	2015-02-13	48.509998	49.049999	47.220001	47.529999	4713100	47.5299
	...	...	...	...	...	...	...
	2015-11-18	27.540001	28.830000	27.309999	28.230000	3091600	28.2300
	2015-11-19	28.190001	28.690001	27.910000	28.059999	1487500	28.0599
	2015-11-20	28.100000	31.250000	28.049999	31.209999	6697500	31.2099
	2015-11-23	30.580000	30.809999	29.150000	29.860001	4029900	29.8600
	2015-11-24	29.459999	30.629999	29.450001	30.010000	2584500	30.0100
	2015-11-25	29.790001	30.540001	29.709999	30.510000	1287100	30.5100
	2015-11-27	30.500000	30.600000	29.610001	30.180000	1058900	30.1800
	2015-11-30	30.110001	30.719999	29.770000	30.129999	2015600	30.1299
	2015-12-01	30.110001	30.459999	29.799999	30.309999	1886000	30.3099
	2015-12-02	30.299999	32.470001	30.290001	31.389999	4650300	31.3899

2015-12-03	31.389999	32.240002	30.480000	30.629999	2698900	30.6299
2015-12-04	30.530001	30.860001	29.320000	30.450001	2313800	30.4500
2015-12-07	30.379999	30.639999	29.629999	30.040001	1362300	30.0400
2015-12-08	29.809999	31.379999	29.500000	30.920000	1830200	30.9200
2015-12-09	30.980000	31.139999	29.260000	30.000000	2238500	30.0000
2015-12-10	30.110001	31.299999	29.990000	30.830000	1252900	30.8300
2015-12-11	30.690001	30.750000	29.600000	29.650000	1415000	29.6500
2015-12-14	29.600000	29.889999	28.850000	29.580000	2328600	29.5800
2015-12-15	29.680000	30.000000	26.459999	26.870001	5759200	26.8700
2015-12-16	26.889999	28.240000	26.260000	28.030001	2992100	28.0300
2015-12-17	28.139999	28.320000	27.190001	27.420000	1483900	27.4200
2015-12-18	27.309999	27.910000	26.900000	27.170000	1299800	27.1700
2015-12-21	27.170000	27.360001	26.030001	26.250000	1947600	26.2500
2015-12-22	26.250000	28.700001	26.150000	27.930000	2952700	27.9300
2015-12-23	27.950001	28.420000	27.440001	28.150000	1001000	28.1500
2015-12-24	28.270000	28.590000	27.900000	28.400000	587400	28.4000
2015-12-28	28.120001	28.379999	27.770000	27.879999	1004500	27.8799
2015-12-29	27.950001	28.540001	27.740000	28.480000	1103900	28.4800
2015-12-30	28.580000	28.780001	28.170000	28.250000	1068000	28.2500
2015-12-31	28.100000	28.969999	28.020000	28.799999	1301500	28.7999

[252 rows x 6 columns]

More on pandas data frames:

<http://pandas.pydata.org/pandas-docs/dev/generated/pandas.DataFrame.html>

## Reading data from a .csv file

```
In [111]: yahoo_stocks.to_csv('yahoo_data.csv')
          #print(open('yahoo_data.csv').read())

In [112]: df = pd.read_csv('yahoo_data.csv')#, index_col='Date', infer_datetime_for
          #print(df.head())
          #print(df.info())
          #print(df.index)
          df
```

```
Out[112]:
```

	Date	Open	High	Low	Close	Volume	\
0	2015-01-02	55.459999	55.599998	54.240002	55.150002	1664500	
1	2015-01-05	54.540001	54.950001	52.330002	52.529999	2023000	
2	2015-01-06	52.549999	53.930000	50.750000	52.439999	3762800	
3	2015-01-07	53.320000	53.750000	51.759998	52.209999	1548200	
4	2015-01-08	52.590000	54.139999	51.759998	53.830002	2015300	
5	2015-01-09	55.959999	56.990002	54.720001	56.070000	6222600	
6	2015-01-12	56.000000	56.060001	53.430000	54.020000	2405100	
7	2015-01-13	54.470001	54.799999	52.520000	53.180000	1952100	
8	2015-01-14	52.799999	53.680000	51.459999	52.200001	1854600	
9	2015-01-15	53.000000	53.610001	50.029999	50.119999	2640400	

10	2015-01-16	50.180000	51.490002	50.029999	51.389999	2183300
11	2015-01-20	51.650002	51.779999	50.689999	51.410000	1227600
12	2015-01-21	51.200001	53.500000	51.200001	53.410000	3248100
13	2015-01-22	53.869999	55.279999	53.119999	54.799999	2295400
14	2015-01-23	54.660000	55.639999	54.299999	55.189999	1636400
15	2015-01-26	55.119999	55.790001	54.830002	55.410000	1450300
16	2015-01-27	56.060001	56.160000	54.570000	55.630001	2410400
17	2015-01-28	56.150002	56.150002	52.919998	53.000000	2013100
18	2015-01-29	52.849998	53.310001	51.410000	52.930000	1844100
19	2015-01-30	52.590000	53.419998	52.049999	52.470001	1875400
20	2015-02-02	52.939999	53.500000	51.209999	53.470001	2105500
21	2015-02-03	53.830002	55.930000	53.410000	55.779999	2885400
22	2015-02-04	55.529999	57.070000	55.250000	56.740002	2498600
23	2015-02-05	57.599998	57.700001	56.080002	57.470001	4657300
24	2015-02-06	47.700001	48.169998	44.860001	45.110001	25137400
25	2015-02-09	44.910000	45.040001	42.099998	42.169998	13079300
26	2015-02-10	43.830002	45.549999	43.310001	44.660000	11267700
27	2015-02-11	45.389999	46.430000	44.810001	46.180000	6359400
28	2015-02-12	46.450001	47.840000	45.950001	47.630001	4375000
29	2015-02-13	48.509998	49.049999	47.220001	47.529999	4713100
..	...	...	...	...	...	...
222	2015-11-18	27.540001	28.830000	27.309999	28.230000	3091600
223	2015-11-19	28.190001	28.690001	27.910000	28.059999	1487500
224	2015-11-20	28.100000	31.250000	28.049999	31.209999	6697500
225	2015-11-23	30.580000	30.809999	29.150000	29.860001	4029900
226	2015-11-24	29.459999	30.629999	29.450001	30.010000	2584500
227	2015-11-25	29.790001	30.540001	29.709999	30.510000	1287100
228	2015-11-27	30.500000	30.600000	29.610001	30.180000	1058900
229	2015-11-30	30.110001	30.719999	29.770000	30.129999	2015600
230	2015-12-01	30.110001	30.459999	29.799999	30.309999	1886000
231	2015-12-02	30.299999	32.470001	30.290001	31.389999	4650300
232	2015-12-03	31.389999	32.240002	30.480000	30.629999	2698900
233	2015-12-04	30.530001	30.860001	29.320000	30.450001	2313800
234	2015-12-07	30.379999	30.639999	29.629999	30.040001	1362300
235	2015-12-08	29.809999	31.379999	29.500000	30.920000	1830200
236	2015-12-09	30.980000	31.139999	29.260000	30.000000	2238500
237	2015-12-10	30.110001	31.299999	29.990000	30.830000	1252900
238	2015-12-11	30.690001	30.750000	29.600000	29.650000	1415000
239	2015-12-14	29.600000	29.889999	28.850000	29.580000	2328600
240	2015-12-15	29.680000	30.000000	26.459999	26.870001	5759200
241	2015-12-16	26.889999	28.240000	26.260000	28.030001	2992100
242	2015-12-17	28.139999	28.320000	27.190001	27.420000	1483900
243	2015-12-18	27.309999	27.910000	26.900000	27.170000	1299800
244	2015-12-21	27.170000	27.360001	26.030001	26.250000	1947600
245	2015-12-22	26.250000	28.700001	26.150000	27.930000	2952700
246	2015-12-23	27.950001	28.420000	27.440001	28.150000	1001000
247	2015-12-24	28.270000	28.590000	27.900000	28.400000	587400
248	2015-12-28	28.120001	28.379999	27.770000	27.879999	1004500

249	2015-12-29	27.950001	28.540001	27.740000	28.480000	1103900
250	2015-12-30	28.580000	28.780001	28.170000	28.250000	1068000
251	2015-12-31	28.100000	28.969999	28.020000	28.799999	1301500

	Adj Close
0	55.150002
1	52.529999
2	52.439999
3	52.209999
4	53.830002
5	56.070000
6	54.020000
7	53.180000
8	52.200001
9	50.119999
10	51.389999
11	51.410000
12	53.410000
13	54.799999
14	55.189999
15	55.410000
16	55.630001
17	53.000000
18	52.930000
19	52.470001
20	53.470001
21	55.779999
22	56.740002
23	57.470001
24	45.110001
25	42.169998
26	44.660000
27	46.180000
28	47.630001
29	47.529999
..	...
222	28.230000
223	28.059999
224	31.209999
225	29.860001
226	30.010000
227	30.510000
228	30.180000
229	30.129999
230	30.309999
231	31.389999
232	30.629999
233	30.450001

```
234 30.040001
235 30.920000
236 30.000000
237 30.830000
238 29.650000
239 29.580000
240 26.870001
241 28.030001
242 27.420000
243 27.170000
244 26.250000
245 27.930000
246 28.150000
247 28.400000
248 27.879999
249 28.480000
250 28.250000
251 28.799999
```

```
[252 rows x 7 columns]
```

The number of tuples in your data

```
In [113]: len(df)
```

```
Out[113]: 252
```

## 1.4 Working with data columns

The columns or “features” in your data

```
In [114]: df.columns
```

```
Out[114]: Index(['Date', 'Open', 'High', 'Low', 'Close', 'Volume', 'Adj Close'], dtype=object)
```

Selecting a single column from your data

```
In [115]: df['Open']
```

```
Out[115]: 0      55.459999
          1      54.540001
          2      52.549999
          3      53.320000
          4      52.590000
          5      55.959999
          6      56.000000
          7      54.470001
          8      52.799999
          9      53.000000
```

10	50.180000
11	51.650002
12	51.200001
13	53.869999
14	54.660000
15	55.119999
16	56.060001
17	56.150002
18	52.849998
19	52.590000
20	52.939999
21	53.830002
22	55.529999
23	57.599998
24	47.700001
25	44.910000
26	43.830002
27	45.389999
28	46.450001
29	48.509998
	...
222	27.540001
223	28.190001
224	28.100000
225	30.580000
226	29.459999
227	29.790001
228	30.500000
229	30.110001
230	30.110001
231	30.299999
232	31.389999
233	30.530001
234	30.379999
235	29.809999
236	30.980000
237	30.110001
238	30.690001
239	29.600000
240	29.680000
241	26.889999
242	28.139999
243	27.309999
244	27.170000
245	26.250000
246	27.950001
247	28.270000
248	28.120001

```
249    27.950001
250    28.580000
251    28.100000
Name: Open, dtype: float64
```

Another way of selecting a single column from your data

```
In [116]: df.Open
```

```
Out[116]: 0      55.459999
1      54.540001
2      52.549999
3      53.320000
4      52.590000
5      55.959999
6      56.000000
7      54.470001
8      52.799999
9      53.000000
10     50.180000
11     51.650002
12     51.200001
13     53.869999
14     54.660000
15     55.119999
16     56.060001
17     56.150002
18     52.849998
19     52.590000
20     52.939999
21     53.830002
22     55.529999
23     57.599998
24     47.700001
25     44.910000
26     43.830002
27     45.389999
28     46.450001
29     48.509998
...
222    27.540001
223    28.190001
224    28.100000
225    30.580000
226    29.459999
227    29.790001
228    30.500000
229    30.110001
```



```

230    30.110001
231    30.299999
232    31.389999
233    30.530001
234    30.379999
235    29.809999
236    30.980000
237    30.110001
238    30.690001
239    29.600000
240    29.680000
241    26.889999
242    28.139999
243    27.309999
244    27.170000
245    26.250000
246    27.950001
247    28.270000
248    28.120001
249    27.950001
250    28.580000
251    28.100000
Name: Open, dtype: float64

```

```
In [117]: df[['Open', 'Close']]
```

```

Out[117]:
      Open  Close
0  55.459999  55.150002
1  54.540001  52.529999
2  52.549999  52.439999
3  53.320000  52.209999
4  52.590000  53.830002
5  55.959999  56.070000
6  56.000000  54.020000
7  54.470001  53.180000
8  52.799999  52.200001
9  53.000000  50.119999
10 50.180000  51.389999
11 51.650002  51.410000
12 51.200001  53.410000
13 53.869999  54.799999
14 54.660000  55.189999
15 55.119999  55.410000
16 56.060001  55.630001
17 56.150002  53.000000
18 52.849998  52.930000
19 52.590000  52.470001
20 52.939999  53.470001

```

21	53.830002	55.779999
22	55.529999	56.740002
23	57.599998	57.470001
24	47.700001	45.110001
25	44.910000	42.169998
26	43.830002	44.660000
27	45.389999	46.180000
28	46.450001	47.630001
29	48.509998	47.529999
...	...	...
222	27.540001	28.230000
223	28.190001	28.059999
224	28.100000	31.209999
225	30.580000	29.860001
226	29.459999	30.010000
227	29.790001	30.510000
228	30.500000	30.180000
229	30.110001	30.129999
230	30.110001	30.309999
231	30.299999	31.389999
232	31.389999	30.629999
233	30.530001	30.450001
234	30.379999	30.040001
235	29.809999	30.920000
236	30.980000	30.000000
237	30.110001	30.830000
238	30.690001	29.650000
239	29.600000	29.580000
240	29.680000	26.870001
241	26.889999	28.030001
242	28.139999	27.420000
243	27.309999	27.170000
244	27.170000	26.250000
245	26.250000	27.930000
246	27.950001	28.150000
247	28.270000	28.400000
248	28.120001	27.879999
249	27.950001	28.480000
250	28.580000	28.250000
251	28.100000	28.799999

[252 rows x 2 columns]

```
In [118]: df.Date.head(10)
```

```
Out[118]: 0    2015-01-02
          1    2015-01-05
          2    2015-01-06
```

```

3      2015-01-07
4      2015-01-08
5      2015-01-09
6      2015-01-12
7      2015-01-13
8      2015-01-14
9      2015-01-15
Name: Date, dtype: object

```

```
In [119]: df.Date.tail(10)
```

```

Out[119]: 242      2015-12-17
          243      2015-12-18
          244      2015-12-21
          245      2015-12-22
          246      2015-12-23
          247      2015-12-24
          248      2015-12-28
          249      2015-12-29
          250      2015-12-30
          251      2015-12-31
Name: Date, dtype: object

```

### Changing the column names

```
In [120]: new_column_names = [x.lower().replace(' ', '_') for x in df.columns]
          df.columns = new_column_names
          df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 252 entries, 0 to 251
Data columns (total 7 columns):
date          252 non-null object
open          252 non-null float64
high          252 non-null float64
low           252 non-null float64
close         252 non-null float64
volume        252 non-null int64
adj_close     252 non-null float64
dtypes: float64(5), int64(1), object(1)
memory usage: 13.9+ KB

```

### Now all columns can be accessed using the dot notation

```
In [121]: df.adj_close
```

```

Out[121]: 0      55.150002
          1      52.529999

```

2	52.439999
3	52.209999
4	53.830002
5	56.070000
6	54.020000
7	53.180000
8	52.200001
9	50.119999
10	51.389999
11	51.410000
12	53.410000
13	54.799999
14	55.189999
15	55.410000
16	55.630001
17	53.000000
18	52.930000
19	52.470001
20	53.470001
21	55.779999
22	56.740002
23	57.470001
24	45.110001
25	42.169998
26	44.660000
27	46.180000
28	47.630001
29	47.529999
	...
222	28.230000
223	28.059999
224	31.209999
225	29.860001
226	30.010000
227	30.510000
228	30.180000
229	30.129999
230	30.309999
231	31.389999
232	30.629999
233	30.450001
234	30.040001
235	30.920000
236	30.000000
237	30.830000
238	29.650000
239	29.580000
240	26.870001

```
241    28.030001
242    27.420000
243    27.170000
244    26.250000
245    27.930000
246    28.150000
247    28.400000
248    27.879999
249    28.480000
250    28.250000
251    28.799999
Name: adj_close, dtype: float64
```

```
In [122]: df.adj_close.head()
```

```
Out[122]: 0    55.150002
          1    52.529999
          2    52.439999
          3    52.209999
          4    53.830002
          Name: adj_close, dtype: float64
```

## 1.5 Data Frame methods

A DataFrame object has many useful methods.

```
In [123]: df.mean()
```

```
Out[123]: open          3.728766e+01
          high          3.805464e+01
          low           3.656373e+01
          close         3.729917e+01
          volume        3.492134e+06
          adj_close     3.729917e+01
          dtype: float64
```

```
In [124]: df.std()
```

```
Out[124]: open          1.128093e+01
          high          1.138111e+01
          low           1.113097e+01
          close         1.125233e+01
          volume        4.145502e+06
          adj_close     1.125233e+01
          dtype: float64
```

```
In [125]: df.median()
```

```
Out [125]: open          3.796500e+01
           high          3.871500e+01
           low           3.637500e+01
           close         3.783500e+01
           volume        2.354050e+06
           adj_close     3.783500e+01
           dtype: float64
```

```
In [126]: df.open.mean()
```

```
Out [126]: 37.287658698412692
```


```
In [127]: df.high.mean()
```

```
Out [127]: 38.054642952380952
```

### 1.5.1 Plotting methods


```
In [128]: df.adj_close.plot(label='adj')
           df.low.plot(label='close')
           plt.legend(loc='best')
```

```
Out [128]: <matplotlib.legend.Legend at 0x11cc669e8>
```



```
In [129]: df.adj_close.hist()
```

```
Out [129]: <matplotlib.axes._subplots.AxesSubplot at 0x11cc704a8>
```



## 1.5.2 Bulk Operations

Methods like **sum()** and **std()** work on entire columns.

We can run our own functions across all values in a column (or row) using **apply()**.

```
In [130]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 252 entries, 0 to 251
Data columns (total 7 columns):
date          252 non-null object
open          252 non-null float64
high          252 non-null float64
low           252 non-null float64
close         252 non-null float64
volume        252 non-null int64
adj_close     252 non-null float64
dtypes: float64(5), int64(1), object(1)
memory usage: 13.9+ KB
```

```
In [131]: df.date.head()
```

```
Out[131]: 0    2015-01-02
          1    2015-01-05
          2    2015-01-06
```

```

3      2015-01-07
4      2015-01-08
Name: date, dtype: object

```

The **values** property of the column returns a list of values for the column. Inspecting the first value reveals that these are strings with a particular format.

```

In [132]: first_date = df.date.values[0]
          first_date

```

```

Out[132]: '2015-01-02'

```

```

In [133]: datetime.strptime(first_date, "%Y-%m-%d")

```

```

Out[133]: datetime.datetime(2015, 1, 2, 0, 0)

```

```

In [134]: df.date = df.date.apply(lambda d: datetime.strptime(d, "%Y-%m-%d"))
          df.date.head()

```

```

Out[134]: 0      2015-01-02
          1      2015-01-05
          2      2015-01-06
          3      2015-01-07
          4      2015-01-08
          Name: date, dtype: datetime64[ns]

```

Each row in a DataFrame is associated with an index, which is a label that uniquely identifies a row.

The row indices so far have been auto-generated by pandas, and are simply integers starting from 0.

From now on we will use dates instead of integers for indices – the benefits of this will show later.

Overwriting the index is as easy as assigning to the index property of the DataFrame.

```

In [135]: df.index = df.date
          df.head()

```

```

Out[135]:
           date      open      high      low      close  volume
date
2015-01-02  2015-01-02  55.459999  55.599998  54.240002  55.150002  166450
2015-01-05  2015-01-05  54.540001  54.950001  52.330002  52.529999  202300
2015-01-06  2015-01-06  52.549999  53.930000  50.750000  52.439999  376280
2015-01-07  2015-01-07  53.320000  53.750000  51.759998  52.209999  154820
2015-01-08  2015-01-08  52.590000  54.139999  51.759998  53.830002  201530

           adj_close
date
2015-01-02  55.150002
2015-01-05  52.529999
2015-01-06  52.439999
2015-01-07  52.209999
2015-01-08  53.830002

```



Now that we have made an index based on date, we can drop the original date column.

```
In [136]: df = df.drop(['date'],axis=1)
          df.info()

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 252 entries, 2015-01-02 to 2015-12-31
Data columns (total 6 columns):
open           252 non-null float64
high           252 non-null float64
low            252 non-null float64
close          252 non-null float64
volume         252 non-null int64
adj_close      252 non-null float64
dtypes: float64(5), int64(1)
memory usage: 13.8 KB
```

### 1.5.3 Accessing rows of the DataFrame

So far we've seen how to access a column of the DataFrame. To access a row we use a different notation.

To access a row by its index value, use the `.ix()` method.

```
In [137]: df.ix[datetime(2015,1,23,0,0)]

Out[137]: open           5.466000e+01
          high           5.564000e+01
          low            5.430000e+01
          close          5.519000e+01
          volume         1.636400e+06
          adj_close      5.519000e+01
          Name: 2015-01-23 00:00:00, dtype: float64
```

To access a row by its sequence number (ie, like an array index), use `.iloc()` ('Integer Location')

```
In [138]: df.iloc[0,:]

Out[138]: open           5.546000e+01
          high           5.560000e+01
          low            5.424000e+01
          close          5.515000e+01
          volume         1.664500e+06
          adj_close      5.515000e+01
          Name: 2015-01-02 00:00:00, dtype: float64
```

To iterate over the rows, use `.iterrows()`

```
In [139]: num_positive_days = 0
         for idx, row in df.iterrows():
             if row.close > row.open:
                 num_positive_days += 1

         print("The total number of positive-gain days is {}".format(num_positive_days))
```

The total number of positive-gain days is 126.

## 1.6 Filtering

It is very easy to select interesting rows from the data.

All these operations below return a new DataFrame, which itself can be treated the same way as all DataFrames we have seen so far.

```
In [140]: positive_days = df[df.close > df.open]
         positive_days.head()
```

```
Out[140]:
```

	open	high	low	close	volume	adj_close
date						
2015-01-08	52.590000	54.139999	51.759998	53.830002	2015300	53.830000
2015-01-09	55.959999	56.990002	54.720001	56.070000	6222600	56.070000
2015-01-16	50.180000	51.490002	50.029999	51.389999	2183300	51.389999
2015-01-21	51.200001	53.500000	51.200001	53.410000	3248100	53.410000
2015-01-22	53.869999	55.279999	53.119999	54.799999	2295400	54.799999

```
In [141]: very_positive_days = df[df.close - df.open > 4]
         very_positive_days.head()
```

```
Out[141]:
```

	open	high	low	close	volume	adj_close
date						
2015-05-07	38.220001	48.73	38.220001	47.009998	33831600	47.009998

```
In [142]: tmp_high = df.high > 60
         tmp_high.head()
```

```
Out[142]:
```

date	high
2015-01-02	False
2015-01-05	False
2015-01-06	False
2015-01-07	False
2015-01-08	False

Name: high, dtype: bool

```
In [143]: len(tmp_high)
```

```
Out[143]: 252
```

Select only the rows of **df** that correspond to **tmp\_high**

```
In [148]: df[tmp_high]
```

```
Out[148]: Empty DataFrame
Columns: [open, high, low, close, volume, adj_close]
Index: []
```

## 1.7 Creating new columns

```
In [149]: df['profit'] = (df.open < df.close)
df.head()
```

```
Out[149]:
```

	open	high	low	close	volume	adj_close
date						
2015-01-02	55.459999	55.599998	54.240002	55.150002	1664500	55.150002
2015-01-05	54.540001	54.950001	52.330002	52.529999	2023000	52.529999
2015-01-06	52.549999	53.930000	50.750000	52.439999	3762800	52.439999
2015-01-07	53.320000	53.750000	51.759998	52.209999	1548200	52.209999
2015-01-08	52.590000	54.139999	51.759998	53.830002	2015300	53.830002

	profit
date	
2015-01-02	False
2015-01-05	False
2015-01-06	False
2015-01-07	False
2015-01-08	True

```
In [150]: for idx, row in df.iterrows():
            if row.close > row.open:
                df.ix[idx, 'gain'] = 'negative'
            elif (row.open - row.close) < 1:
                df.ix[idx, 'gain'] = 'small_gain'
            elif (row.open - row.close) < 6:
                df.ix[idx, 'gain'] = 'medium_gain'
            else:
                df.ix[idx, 'gain'] = 'large_gain'
df.head()
```

```
Out[150]:
```

	open	high	low	close	volume	adj_close
date						
2015-01-02	55.459999	55.599998	54.240002	55.150002	1664500	55.150002
2015-01-05	54.540001	54.950001	52.330002	52.529999	2023000	52.529999
2015-01-06	52.549999	53.930000	50.750000	52.439999	3762800	52.439999
2015-01-07	53.320000	53.750000	51.759998	52.209999	1548200	52.209999
2015-01-08	52.590000	54.139999	51.759998	53.830002	2015300	53.830002

	profit	gain
date		
2015-01-02	False	small_gain

```

2015-01-05  False  medium_gain
2015-01-06  False  small_gain
2015-01-07  False  medium_gain
2015-01-08   True   negative

```

## Alternatively

```

In [ ]: def namerow(row):
        if row.close > row.open:
            return 'negative'
        elif (row.open - row.close) < 1:
            return 'small_gain'
        elif (row.open - row.close) < 6:
            return 'medium_gain'
        else:
            return 'large_gain'

df['test_column'] = df.apply(namerow, axis = 1)

In [ ]: print df.head()

In [ ]: df.drop('test_column', axis = 1)

```

## 1.8 Grouping

Besides `apply()`, another great DataFrame function is `groupby()`. It will group a DataFrame by one or more columns, and let you iterate through each group.

Here we will show the average gain among the three categories of negative, small, medium and large gains we defined above and stored in column “gain”

```

In [ ]: gains = {}
        for gain, gain_data in df.groupby("gain"):
            gains[gain] = gain_data.close.mean()
        gains

```

## 1.9 Comparing multiple stocks

As a last task, we will use the experience we obtained so far – and learn some new things – in order to compare the performance of different stocks we obtained from Yahoo finance.

```

In [ ]: stocks = ['ORCL', 'TSLA', 'IBM', 'YELP', 'MSFT']
        attr = 'Close'
        df = pd.io.data.get_data_yahoo(stocks,
                                         start=datetime(2014, 1, 1),
                                         end=datetime(2014, 12, 31))[attr]

        print df.head()

In [ ]: df.ORCL.plot(label = 'oracle')
        df.TSLA.plot(label = 'tesla')

```

```
df.IBM.plot(label = 'ibm')
df.MSFT.plot(label = 'msft')
df.YELP.plot(label = 'yelp')
plt.legend(loc=5)
```

### Calculating returns over a period of length T

$$r(t) = \frac{f(t) - f(t - T)}{f(t)}$$

The returns can be computed in python with a simple function `pct_returns()`

```
In [ ]: rets = df.pct_change(30)
        #print rets.head(20)
```

Plotting again the timeseries of the returns of the different stocks

```
In [ ]: rets.ORCL.plot(label = 'oracle')
        rets.TSLA.plot(label = 'tesla')
        rets.IBM.plot(label = 'ibm')
        rets.MSFT.plot(label = 'msft')
        rets.YELP.plot(label = 'yelp')
        plt.legend()

In [ ]: plt.scatter(rets.TSLA, rets.ORCL)
        plt.xlabel('Returns TESLA')
        plt.ylabel('Returns ORCL')
```

**Correlations of columns** The correlation coefficient between variables  $X$  and  $Y$  is defined as follows:

$$\text{Corr}(X, Y) = \frac{E[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}$$

In python we can compute the correlation coefficient of all pairs of columns with `corr()`

```
In [ ]: corr = rets.corr()
        print corr
```

**Visualizing the correlation coefficient of all columns** We will learn more about visualization later, but for now this is a simple example

```
In [ ]: sns.heatmap(corr, annot=True)
```

**Returns vs risk (standard deviation)** In many applications, we want to know both the returns as well as the standard deviation of the returns of a stock (i.e., its risk). Below we visualize the result of such an analysis

```

In [ ]: plt.scatter(rets.mean(), rets.std())
        plt.xlabel('Expected returns')
        plt.ylabel('Standard Deviation (Risk)')
        for label, x, y in zip(rets.columns, rets.mean(), rets.std()):
            plt.annotate(
                label,
                xy = (x, y), xytext = (20, -20),
                textcoords = 'offset points', ha = 'right', va = 'bottom',
                bbox = dict(boxstyle = 'round,pad=0.5', fc = 'yellow', alpha = 0.5),
                arrowprops = dict(arrowstyle = '->', connectionstyle = 'arc3,rad=0')

In [ ]: # Code for setting the style of the notebook
        from IPython.core.display import HTML
        def css_styling():
            styles = open("../theme/custom.css", "r").read()
            return HTML(styles)
        css_styling()

In [ ]:

```